

Artem Yushkovskiy

AI PLATFORM ENGINEER & TECH LEAD • DISTRIBUTED AI ORCHESTRATION • CLOUD-NATIVE (K8S, GCP, AWS)

Berlin, Germany · ajuszkowski@gmail.com · [atemate](#) · [atemate](#)

Summary

AI Platform Engineer and Tech Lead with 10+ years in software engineering. Speaker at KubeCon EU 2026 and creator of an open-source actor mesh for orchestrating AI workloads on Kubernetes. Operating at the intersection of data science, backend engineering, DevOps, and MLOps - building modular, production-grade platforms grounded in stateless execution, durable workflows, and linear scalability.

Highlights

- ◆ Built self-hosted AI image processing platform on AWS/Kubernetes processing millions of images at near-realtime throughput - AWS EKS, self-hosted Stable Diffusion models, GPU auto-scaling (KEDA), linear horizontal autoscaling
- ◆ Created [asya.sh](#) 🤖 - Open-source Kubernetes-native actor mesh for async AI orchestration - Apache 2.0; designed as foundation for AI ops platforms
- ◆ Driving company-wide initiatives: self-hosted AI adoption, GPU cost reduction, cross-team dataset versioning standard
- ◆ Advanced agentic software engineering - coordinating tens of parallel AI agents across large production codebases; author of [git-aint](#); actively transferring expertise across teams
- ◆ Google Cloud Certified: Professional Architect (2025) + Professional ML Engineer (2023)

Experience

Delivery Hero – Berlin, DE

Aug 2023 – present

Senior AI Engineer - Distributed AI Orchestration, GenAI

- ◆ Designed and scaled a self-hosted AI image processing platform on AWS/Kubernetes (outpainting, upscaling, watermark/object removal), which evolved from batch ML pipelines to a dedicated GPU cluster to an actor-based async inference architecture; processing millions of images at near-realtime throughput; working closely with DS on model experimentation, transparency, and testability
- ◆ Designed CI/CD for monorepo with 20+ model packages building multi-GB GPU images enabling O(1) effort when onboarding new models; single parameterized Dockerfile; build time 4h → 20 min (cache miss), 1h → 6 min (cache hit); pod startup 7 min → 3 min; image sizes 3–4× smaller [Docker Buildx, uv, optimized layer caching]
- ◆ Actor mesh architecture proved highly effective for generic AI use-cases beyond image processing; open-sourced as [asya.sh](#) 🤖; driving internal adoption across teams with growing interest from the external community
- ◆ Authored company-wide RFC on GPU capacity and cost strategy (training + serving) covering 18-36 month demand forecasting, usage patterns across teams, and negotiation baseline for cloud provider agreements; designed cross-team dataset versioning standard enabling model comparability across dataset evolution
- ◆ Mentoring junior engineers on clean code and architecture design; coaching data scientists on software engineering best practices to raise code quality and testability across the team
- ◆ Stack: *Python asyncio, Go, Kubernetes* (CRDs, Helm, Kustomize), *KEDA, ArgoCD, FluxCD, Drone CI, GitHub Actions, OpenTelemetry, Grafana, Terraform, AWS*

Delivery Hero – Berlin, DE

June 2022 – Aug 2023

ML Engineer - Incentive Fraud Detection

- ◆ Owned full lifecycle of 20+ low-latency production ML models for voucher fraud detection across 70+ countries: Vertex AI training pipelines, batch inference, feature stores (BigQuery, Airflow, DynamoDB), and serving optimized to <10ms latency on every food order request
- ◆ enabling realtime blocking of incentive fraud at scale, saving hundreds of thousands monthly
- ◆ Built end-to-end CI/CD for ML models integrating GitHub Actions with Vertex AI Pipelines and Cloud Build
- ◆ enabling rapid model iteration while maintaining MLOps best practices (dataset versioning, model registry, automated evaluation)
- ◆ Engineered a dataset versioning approach to ensure strict model comparability; later recognized this as a systemic bottleneck and successfully pitched it as a platform-wide standard
- ◆ On-call SRE duties, monitoring and observability (DataDog, OpsGenie)
- ◆ Co-authored with Google: "[How Delivery Hero connected GitHub with Vertex AI to manage 20+ voucher fraud detection models](#)" (Google Cloud Blog, Nov 2023)
- ◆ Stack: GCP Vertex AI, BigQuery, Airflow, DynamoDB, Terraform, AWS EKS, GitHub Actions, FastAPI, DataDog

Delivery Hero – Berlin, DE

Nov 2021 – June 2022

ML Engineer - ML Platform

- ◆ Designed and built an internal ML Platform on GCP Vertex AI
- ◆ providing shared data and ML pipeline infrastructure across teams (CLV, dish normalisation, email suspicion, fraud detection)
- ◆ Enabled consistent model training, evaluation, and deployment workflows across multiple AI use-cases and teams

Neuromation / Neu.ro – St Petersburg, RU

Jan 2018 – Jan 2021

Software Engineer → MLOps Engineer

- ◆ Built core of a Kubernetes-based ML platform ("Docker-like experience on any cloud and bare metal"): microservices, CLI/SDK, resource orchestration; later formed and joined MLOps tooling team
- ◆ full ML lifecycle, client support, Kubernetes internals (Prometheus, Grafana, DVC, Seldon, Feast)

Positive Technologies + Aalto University

Jan 2015 – Jan 2018

Early career

- ◆ Static code analysis, security vulnerability research, distributed systems thesis

Projects

Asya 🦄 - Kubernetes-native Actor Mesh for AI Orchestration'

Nov 2025

- ◆ Actor-based framework for building async, decentralized AI systems on Kubernetes
- ◆ Pure Python handlers + CRD-driven infrastructure: clean separation of business logic and platform configuration
- ◆ Supports A2A protocols, dynamic routing, streaming, and human-in-the-loop flows
- ◆ Built-in resilience via actor flavors (CRD overlays): retry policies, timeouts, dead-letter queues
- ◆ defined as deployment config, not code
- ◆ KEDA-based independent per-actor autoscaling including scale-to-zero
- ◆ Presented at KubeCon + CloudNativeCon EU 2026
- ◆ [asya.sh](#) | [GitHub](#)

git-aint - AI-native Issue Tracker

Mar 2026

- ◆ Issue tracking where tasks are markdown files stored in git (separate worktree)
- ◆ no external database, no SaaS dependency
- ◆ Designed for parallel multi-agent development workflows while keeping human in the loop: AI agents read and write tasks natively alongside humans, all stored in plain text in git
- ◆ Enables RFC, ADR, tech design, and task tracking as editable .md files
- ◆ complexity owned by the team, not hidden in a third-party tool
- ◆ Together with git worktrees, allows parallelizing implementation across tens of agents without losing control on the code
- ◆ [GitHub](#)

Public Speaking

- ◆ "REST in Peace: AI Needs to Be Async" - KubeCon + CloudNativeCon EU 2026 (~13k attendees)
- ◆ AI Infrastructure Summit Europe 2026, Munich - we.CONECT (500+ attendees)
- ◆ "Fast & Asynchronous: Drift Your AI, Not Your GPU Bill" - AI in Production Meetup, Berlin (150+ participants, GetYourGuide, Nov 2025)
- ◆ "Fast & Asynchronous: Drift Your AI, Not Your GPU Bill" - Agents in Production virtual event (~4k participants, MLOps Community x Prosus, Dec 2025)
- ◆ organized regular meetups for MLOps Community Berlin; was leading Ask-Me-Anything initiative in MLOps Community Slack (2021-2023)

Education

Aalto University (Helsinki, FI) & ITMO University (St Petersburg, RU)

Jan 2016 – Jan 2018

MSc in Computer Science

- ◆ Double-degree programme: Information Security and Cloud Computing, with honors. Thesis: Automated Analysis of Weak Memory Models.

ITMO University (St Petersburg, RU)

Jan 2012 – Jan 2016

BSc in Computer Science

- ◆ Information Security, with honors.

Skills

AI Systems & Inference: self-hosted model serving, GPU orchestration (KEDA, scale-to-zero), async actor-based pipelines, AI guardrails, image generation (SDXL), LLM integration, vLLM, Triton, PyTorch, HuggingFace, ONNX

Platform Engineering: Kubernetes (CRDs, operators, Helm, Kustomize), GitOps (ArgoCD, Flux), Terraform, CI/CD (GitHub Actions, Drone CI, Docker Buildx, Kaniko), observability (OpenTelemetry, Prometheus, Loki, Grafana, DataDog), alerting (OpsGenie), secret management (Vault)

MLOps & Data: ML lifecycle (Vertex AI, MLflow, W&B), pipeline orchestration (Airflow, KFP, Flyte, Metaflow), feature stores, dataset versioning (DVC), data quality (Great Expectations, Evidently), stream processing (Flink), low-latency model serving (<10ms)

Languages & Infrastructure: Python (asyncio, FastAPI, uv), Go, Bash, Linux, AWS (EKS, SQS, SNS, Lambda, ECR), GCP (Vertex AI, Pub/Sub, Cloud Run, Cloud Build, GCS, Artifact Registry), PostgreSQL, DynamoDB, Redis, RabbitMQ, NATS, MongoDB

Agentic Engineering: multi-agent parallel development, agentic frameworks (A2A, MCP, LangGraph, CrewAI), Langfuse, Langwatch, Scenarios, Claude Code, beads, git-aint

Languages

Fluent: English, Russian

Intermediate: French

Basic: German

Certifications

- ◆ Google Cloud Professional Architect (2025)
- ◆ Google Cloud Professional ML Engineer (2023)